

Fareez Ahmed

Staff Full Stack & Applied AI Engineer | **Active Secret Clearance**

San Francisco, CA / Dallas, TX (Remote)

(202) 677-0160 · Fareez.Ahmed@gmail.com

v2software.ai · linkedin.com/in/fareezahmed

github.com/youngfreezy · huggingface.co/fareezaiahmed

Software engineer with ten years of web and backend experience, the last few on production LLM systems: agent runtimes on LangGraph and Google ADK, A2A and MCP integration, model gateways with routing and cost controls, and evals before rollout. Now also on the evaluation side: an AI Fellow at Handshake AI writing coding tasks that frontier models fail and grading agent runs, and an evaluator of model-written documents for OpenAI through Mercor. Owns delivery across React and TypeScript interfaces, Python/FastAPI services, and AWS infrastructure. Built Mayo Clinic's article-validation workflow, cutting review time from about two hours to under ten minutes, and delivered Signet Jewelers' internal AI agent platform in three months.

TECHNICAL SKILLS

AI systems	LangGraph, Google ADK, A2A, MCP servers and clients, LangChain, LlamaIndex, Amazon Bedrock and AgentCore, Anthropic and OpenAI APIs, RAG, structured outputs and tool calling, Inngest
Model routing	LiteLLM, OpenRouter, Bedrock and Databricks AI Gateway, tiered routing with escalation and cross-provider failover, per-team budgets, rate limits, and cost attribution
LLM evals and safety	RAGAS, MLflow GenAI, LLM-as-judge, rubric design and calibration, behavioral regression evals in CI, human-in-the-loop approval gates, kill switches, redacted tracing
Frontend	React, Next.js, TypeScript, JavaScript, Redux, Tailwind, Angular, D3.js
Backend	Python, FastAPI, Flask, Node.js, Express, NestJS, Java, Ruby on Rails
Databricks	Jobs, Unity Catalog, Delta Lake, Asset Bundles, Model Serving, AI Gateway, inference tables, Lakebase Postgres, MLflow 3
Cloud and delivery	AWS Lambda, S3, EC2, ECS Fargate, EKS, Step Functions, SQS, CDK, serverless; Terraform, Terragrunt, Kubernetes, ArgoCD, Istio; Vertex AI, Azure ML, GitHub Actions, GitLab CI
Data and testing	PostgreSQL, PGVector, Neo4j, DynamoDB, MongoDB, Redis, ingestion pipelines, SQL migrations, OpenTelemetry; Playwright, Cypress, Jest, Pytest, Gherkin BDD, React Testing Library, Storybook

AI MODEL EVALUATION AND TRAINING DATA

AI Fellow, Handshake AI

Jul 2026 to present

Model evaluation and training tasks for frontier AI labs, including Meta:

- Project Dynamo: write coding tasks that frontier models cannot yet solve. Each task is a Docker environment with a hidden verifier and a reference solution. A task is accepted only when the lab's agent fails most of its trials.
- Built a LangGraph toolkit that runs the Dynamo pipeline as a resumable state machine with approval gates. It checks a task locally before it costs a CI run, attacks my own verifier to find holes, and calibrates difficulty against real trial results. About 8,500 lines of Python across 20 commands.
- Project Roadhouse (Meta): compare two coding agents on one staged task in isolated Docker containers, across Claude Code, Codex, OpenCode, Kimi Code and Meta's Muse Code. Tasks span Python asyncio and TLS, DuckDB internals in C++, Rust, U-Boot in C, TypeScript, Hack/HHVM and a PyTorch-to-JAX port.
- Verify each stage myself: rerun benchmarks and load tests (a 50,000-request HTTP flood) and write my own probes; on an asyncio TLS task, two probes found a different regression in each agent's patch.
- Rank each pair on correctness, agent behavior and communication, citing turns, test results and hashed evidence.

AI Evaluator, OpenAI via Mercor

Aug 2026 to present

- Evaluated hundreds of model-written business documents and slide decks on calibrated multi-dimension writing and visual-quality rubrics, each with a documented rationale.
- Built rubric-conformance checks and two independent fresh-context AI judges that review every draft before submission.

EXPERIENCE

Staff Software Engineer, V2 Software LLC

Nov 2018 to present

Selected client engagements and engineering contributions:

Polaris IO *Applied AI platform*

- Lead engineering across Python/FastAPI services, a TypeScript/Node backend, React interfaces, Playwright tests, and AWS EKS deployments.
- Created the agent SDK and LangGraph runtime: services request a model tier rather than a specific model, calls escalate to a stronger tier when output fails schema validation, fail over across providers during outages, and are checked against budget before they run. Versioned contracts keep the Python and TypeScript services aligned. Chose LangGraph through a 144-run comparison of candidate frameworks.
- Delivered client-level LLM cost attribution, budget caps, and kill switches, with recorded usage reconciled against OpenAI's and Anthropic's own usage reports. Verified enforcement live against both providers.
- Built encrypted storage and replay of completed agent runs in tenant-scoped PostgreSQL; engineers inspect saved model and tool responses without repeating calls or writes.
- Built the observability pipeline: Datadog tracing with sensitive-field redaction that blocks unredacted events from being logged.

- Built PioWorkflow, the team's AI developer toolkit: QA where a judge model must cite evidence before work passes, an evaluation framework that cut agent-skill testing cost from \$4.49 to under \$1 per run, and a session registry that coordinates worktrees and file ownership across coding agents (POSIX and Windows CI).

eSimplicity *Rapid Innovation Center*

- Built the Astro/React frontend for SHINE, an AI chat and search portal on Databricks for government teams, designed as a reference system other agencies clone and run with configuration changes. Added chat history, role-based access, per-user rate limits, and reusable themes built to WCAG targets.
- Provisioned Databricks model serving and AI gateway infrastructure as code, connected the chat agent to SharePoint documents, and implemented authenticated automated deployments.
- Co-built Katalyst Forge, an AI software-delivery platform: consolidated 40 agents into a versioned release, routed design, review, and QA stages to frontier models and routine code generation to a cheaper Bedrock model under spend ceilings, and added cost tracking, evaluation gates, and A2A support for ADK agents.
- Delivered two internal products independently: a bot that publishes a checked visual guide to each merged change (500+ brief library, 333+ authored by me), and a Medicaid data-quality demo with an automated behavioral test suite and release controls requiring verifiable human approval.
- Deployed shared Kubernetes services: feature flags with single sign-on, artifact-integrity checks for the container supply chain, and a pull-request review bot using a model hosted inside the tenant.

Mayo Clinic *AI content validation*

- Built a five-agent workflow that cut health-article review from about two hours to under ten minutes: agents check an article in parallel behind a router that clears straightforward cases without a model call.
- Orchestrated the workflow in LangGraph, using interrupt() so no result reached an editor before a person reviewed it. Built the checking agents on Google ADK and connected them over A2A, so each runs as its own service with a published agent card.
- Built the loader for 34 validation rules: graph database at runtime with a file fallback, so a database outage degrades to a slower path instead of downtime. Measured groundedness and coverage with RAGAS before rollout.
- Built a Next.js reviewer dashboard with streaming results and a human approval step before publication; developed React components and auth and logging middleware for the content platform.

Signet Jewelers *AI Foundry*

- Delivered AI Foundry in three months: an AWS-based studio where internal teams build agents that search their own documents.
- Built an Amazon Bedrock gateway as the single entry point for model traffic, with per-team budgets, rate limits, response caching, and token metrics by team, model, and project; rerouting or shutting off a model was a configuration change, not a deploy.
- Built streaming chat over Bedrock with keep-alive heartbeats so long responses survive CDN idle timeouts.
- Built a three-agent code-review copilot (planning, orchestration, QA) that cut pull-request review from days to hours.

JPMorgan Chase

- React/Redux microfrontends for portfolio and lending workflows, Node services on AWS serverless infrastructure, and partner integrations and analytics pipelines.

CACI / Army

- Led React UI delivery on secure DoD networks and built Python services for data processing and ML-output visualization.

OPEN SOURCE

Author of **JobHunter Agent** (github.com/youngfreezy/job-hunter-agent, live at jobhunteragent.com), an open-source autonomous job-application platform: an eight-agent LangGraph pipeline behind FastAPI and Next.js, with Bright Data MCP job discovery, Skyvern form filling, human-approval gates before any submission, SSE streaming, encrypted resume storage, and Stripe billing.

Author of **unslopify**, a Python writing-quality linter on PyPI. Implements 27 deterministic checks across four categories and a cross-document phrase bank; can run in CI without model calls.

ADDITIONAL CLIENT WORK

Bayer Radiology (medical imaging): Angular 9, TypeScript, NgRx, and Electron desktop software for radiology imaging and catheter-based contrast injection procedures; built new features and refactored legacy code on a Node.js, MongoDB, and Socket.IO stack.

Rivian: DynamoDB and Python Lambda supply-chain pipelines. **College Board:** React microfrontends and AWS serverless services for Student Portal financial aid, used by millions of students. **Live Nation / Ticketmaster:** React, Next.js, and GraphQL for event discovery and ticketing. **Etsy / Reverb:** Led Plaid integration and built reusable React payment components. **Department of Veterans Affairs:** React and Rails internationalization and CMS work.

EARLIER ENGINEERING ROLES

Senior Software Engineer, CoStar Realty Group. Express mobile APIs, React listings, AWS Kinesis pipelines. Jun 2017 to Nov 2018

Senior Software Engineer, DataScan. Angular, D3, and Node/Express wholesale-loan platform. Aug 2016 to Mar 2017

Software Engineer, General Electric. Built supply-chain software in AngularJS and Node.js. Nov 2015 to Aug 2016

EDUCATION & CERTIFICATION

MA, International Affairs and Economics, Columbia University. GPA 3.62.

BA, Economics, George Washington University. Summa cum laude, Phi Beta Kappa, GPA 3.8.

Agentic AI Engineering, DeepLearning.AI, Oct 2025.

Languages: American English, French, Arabic, Bengali.